



```

const res = await client.query(
  'SELECT id, url, description, created_at from links')
result = res.rows
await client.end()
} else if(event.method == "POST") {
  const client = new Client(opts)
  await client.connect()

  const res = await client.query(
    'INSERT INTO links(id, url, description, created_at)' +
    ' VALUES (DEFAULT, $1, $2, now())',
    [event.body.url,
     event.body.description])
  result = res.rows[0]
  await client.end()
}

return context
  .status(200)
  .succeed(result)
}

```

Create the initial schema with psql, you may need to install the following package to get the CLI tool:

```
sudo apt install postgresql-client
```

```
psql -h 10.62.0.1 -U postgres
Password for user postgres:
```

Then paste in the schema from above.

```

\dt
      List of relations
 Schema | Name  | Type  | Owner
-----+-----+-----+-----
 public | links | table | postgres
(2 rows)

SELECT * from links;
 id | created_at | url | description
---+-----+-----+-----
(0 rows)

```

Deploy the function:

```
faas-cli up -f push-pull.yml
```

Test the function by posting a link:

```

echo '{ "url": "https://inlets.dev",
  "description": "Cloud Native Tunnels"
}' |
curl --data-binary @- \

```

```
--header "Content-Type:application/json" \
$OPENFAAS_URL/function/push-pull
```

Check it was inserted:

```
curl -s $OPENFAAS_URL/function/push-pull
```

```
{
  "id": 3,
  "url": "https://inlets.dev",
  "description": "Cloud Native Tunnels",
  "created_at": "2021-01-18T13:26:04.187Z"
}
```

Hint: If you have the `jq` utility installed, then you can pipe the output to `jq` to pretty-print it.

You can now store and query data from a relational database. There are more efficient ways of accessing a database than creating a connection per request, which you can learn about in the [postgres-node docs](#) under the "pool" section.

The `node12` template reuses the same memory for each request, meaning you can obtain a pool of connections once and re-use it for each request. This is beyond the scope of the eBook.

Bonus task: Given what you've learnt about using API tokens in the previous, can you update the `push-pull` function to only allow POSTs when a valid token is passed?

Unit testing with Node.js

When using an OpenFaaS template, the `Dockerfile` can optionally specify a command to run unit tests during the `faas-cli build/publish`.

Look at the `template/node12/Dockerfile` and you will run a line that reads: `npm test`.

You can unit test your function's handler or any library functions that you may write. The following example uses the [mocha](#) test runner, and the [chai](#) assertion library to help you verify the results.

We'll use a Test-Driven Development style to:

- Write a function that returns a 200 status code
- Then write a unit test that expects a 201 status code
- See it fail
- Then update the code to 201 in the `handler.js`
- And see it pass

Create a new function:

```
export OPENFAAS_PREFIX="alexellis2"
```

```
faas-cli new --lang node12 api
```

Edit `handler.js`:

```
'use strict'

module.exports = async (event, context) => {
  const result = {
    "status": "ok"
  }
```

```

    }

    return context
      .status(200)
      .succeed(result)
  }
}

```

Install the mocha and chai npm modules:

```

cd api
npm install --save mocha chai

```

These will now be listed in package.json.

Next, modify the test step in the package.json file:

```

"scripts": {
  "test": "mocha test.js"
}

```

Now create your first test:

```

'use strict'

const expect = require('chai').expect
const handler = require('./handler') // The function's handler

describe('Our API', async function() {
  it('gives a 201 for any request', async function() {
    let cb = async function (err, val) { };
    let context = new FunctionContext(cb);
    let event = new FunctionEvent({body: ''})
    let res = await handler(event, context)

    expect(context.status{}).to.equal(201)
  });
});

```

Then, add the following test-doubles underneath for the FunctionEvent and FunctionContext. These test-doubles (also known as stubs) allow you to provide canned values for the input, and to inspect the results.

```

class FunctionEvent {
  constructor(req) {
    this.body = req.body;
    this.headers = req.headers;
    this.method = req.method;
    this.query = req.query;
    this.path = req.path;
  }
}

class FunctionContext {
  constructor(cb) {
    this.statusValue = 200;
    this.cb = cb;
  }
}

```

```

    this.headerValues = {};
    this.cbCalled = 0;
  }
  status(value) {
    if(!value) {
      return this.statusValue;
    }

    this.statusValue = value;
    return this;
  }
  headers(value) {
    if(!value) {
      return this.headerValues;
    }

    this.headerValues = value;
    return this;
  }
  succeed(value) {
    let err;
    this.cbCalled++;
    this.data = value;
    this.cb(err, value);
  }
  fail(value) {
    let message;
    this.cbCalled++;
    this.cb(value, message);
  }
}

```

You can find the whole file in this GitHub repo: [alexellis/openfaas-node12-mocha-unit-test/](https://github.com/alexellis/openfaas-node12-mocha-unit-test/)

Now you can `faas-cli build`, and you will see mocha running and printing results to the console. If the tests fail, then mocha will produce a non-zero exit code and the build will stop.

```

> openfaas-function@1.0.0 test /home/app/function
> mocha test.js

```

```

Our API
  1) gives a 201 for any request

0 passing (8ms)
1 failing

1) Our API
   gives a 201 for any request:

AssertionError: expected 200 to equal 201
+ expected - actual

```

```
-200
+201

at Context.<anonymous> (test.js:13:37)
```

You can also run the unit-tests on your local machine if you have Node.js and npm already installed:

```
cd api/
npm test
```

Now edit the handler to return a 201 status so that the unit test passes.

Then run `faas-cli build` again.

Here is an example of the unit test passing:

```
---> Running in 47a902a0d06a

> openfaas-function@1.0.0 test /home/app/function
> mocha test.js

  Our API
    [x] gives a 201 for any request

  1 passing (6ms)
```

Removing intermediate container 47a902a0d06a

You can also customise the test-doubles for `FunctionEvent` and `FunctionContext` or create your own npm module for them, to reduce duplication. The same technique should apply to other unit-test runners such as [Jest](#), which is more popular with the [React](#) community.

How to build multiple functions

You can have multiple OpenFaaS stack files with a single function in them, but it makes sense to use a single file to filter it when required. It simplifies maintenance and CI/CD.

Hint: If you rename your function's YAML file to "stack.yml", you no longer need to use the `-f` parameter.

You can append additional functions into your stack file. Then you can build and deploy them in one shot, or filter down to just the one which has changes:

```
faas-cli new --lang node12 hackernewsbot --append stack.yml
```

We'll now have both `starbot` and `hackernewsbot` in our YAML file with their own handlers and own package.json files.

```
version: 1.0
provider:
  name: openfaas

functions:
  starbot:
    lang: node12
    handler: ./starbot
```

```
image: alexellis2/starbot:latest

hackernewsbot:
  lang: node12
  handler: ./hackernewsbot
  image: alexellis2/hackernewsbot:latest
```

If you run any of the commands like build/publish/up or deploy, then everything in the file will be used by default. To filter to just one use `--filter NAME`.

For example:

```
faas-cli up --filter hackernewsbot
```

A super-power of `faas-cli` is how it can enable parallel builds of functions using Docker. Here's an example to build up to two functions at once:

```
faas-cli up --parallel 2
```

Invoking functions

Functions can be invoked either synchronously (the caller waits until it's done), or asynchronously (the caller gets an `X-Call-ID` header back and doesn't have to wait).

Typically, you will want to use a synchronous invocation, which is the simplest and easiest to consume.

You can find your function's URL using `faas-cli describe` or by opening the OpenFaaS UI dashboard at <http://127.0.0.1:8080>

```
faas-cli describe figlet
Name:          figlet
Status:        Ready
Replicas:      1
Available replicas: 1
Invocations:   0
Image:
Function process:
URL:           http://127.0.0.1:8080/function/figlet
Async URL:     http://127.0.0.1:8080/async-function/figlet
```

The simplest way to invoke your function is by using its URL exposed by the gateway: <http://127.0.0.1:8080/function/figlet>

You can also invoke the function via the CLI:

```
echo faasd | faas-cli invoke figlet
```

```

_ _ _ _ _
|_| / _ \ / _ \ / _ \ / _ \
|_| ( ) ( ) ( ) ( ) ( ) ( )
|_| \_/_/_/_/_/_/_/_/_/_/_
```

Now, if you're doing something like generating PDFs, and you are either processing a lot of them, or they take a while to produce, then you will want to consider asynchronous invocations. This is the same as using a queue, except that the result is thrown away by default.

To receive the result, simply pass in a *Callback URL*. You will be able to know which result is for which invocation by matching on the *Call ID*.

```
curl --data-binary input.jpg http://127.0.0.1:8080/async-function/resize/?width=500px \
  --header "X-Callback-Url: http://example.com/api/callback"
```

The `faas-cli` can also be used to invoke functions asynchronously, with a slightly different syntax:

```
cat input.jpg | faas-cli invoke --async analyse-image \
  --header "X-Callback-Url=http://example.com/api/callback"
```

The `queue-worker` will then send you the result and a HTTP code (depending on success or failure) to `http://example.com/api/callback`

The recipient of the callback could be another function deployed with `faasd`, or any other HTTP endpoint.

To match a request with a response use the `X-Call-Id` header received when you invoked the function.

Making any process into a function

A common use-case we have seen with OpenFaaS is to take a CLI and make it into a function. You can do this with anything that can be installed into a container, from ImageMagick to the AWS CLI to `ffmpeg`.

There are two good blog posts for this:

- [Turn Any CLI into a Function with OpenFaaS](#)
- [Stop installing CLI tools on your build server — CLI-as-a-Function with OpenFaaS](#)

The easiest example is to use the "dockerfile" template and have it run a pre-installed bash command:

```
# Customise this to your own username
export OPENFAAS_PREFIX=alexellis2
faas-cli new --lang dockerfile env
```

The `fprocess` line tells the container what to run as the function, change it to `env`:

```
FROM ghcr.io/openfaas/classic-watchdog:0.1.4 as watchdog

FROM alpine:3.12

RUN mkdir -p /home/app
COPY --from=watchdog /fwatchdog /usr/bin/fwatchdog
RUN chmod +x /usr/bin/fwatchdog

# Add non root user
RUN addgroup -S app && adduser app -S -G app
RUN chown app /home/app
WORKDIR /home/app

USER app

# Populate example here - i.e. "cat", "sha512sum" or "node index.js"
ENV fprocess="env"

CMD ["fwatchdog"]
```

Now deploy it, and remember that Raspberry Pi users need to run `faas-cli publish/deploy` instead.


```
faas-cli up -f env.yml
```

When you invoke it, you'll see all the HTTP headers injected as environment variables:

```
echo | faas-cli invoke env
```

```
PATH=/usr/local/sbin:/usr/local/bin:/usr/sbin:/usr/bin:/sbin:/bin
fprocess=env
HOME=/home/app
Http_X_Forwarded_Host=faasd.alex.o6s.io
Http_User_Agent=curl/7.68.0
Http_Content_Type=application/x-www-form-urlencoded
Http_X_Forwarded_For=81.99.136.188
Http_Accept-Encoding=gzip
Http_X_Forwarded_Proto=https
Http_Content_Length=0
Http_Accept=/*/*
Http_Method=POST
Http_ContentLength=0
Http_Path=/
Http_Host=10.62.0.92:8080
```

Let me show you how to run curl from inside a function. This is very useful for testing connectivity and DNS.

Now add this line to your Dockerfile before USER app to add the curl package into the container.

```
RUN apk add --no-cache curl
```

Now change the fprocess to xargs curl in the Dockerfile:

```
ENV fprocess="xargs curl"
```

Run faas-cli up -f env.yml again to redeploy the function.

Now test it out:

```
curl -sLs https://faasd.example.com/function/env \
--data "-s http://api.open-notify.org/astros.json"
```

Abbreviated output:

```
{
  "message": "success",
  "number": 3,
  "people": [
    {
      "craft": "ISS",
      "name": "Sergey Ryzhikov"
    },
    {
      "craft": "ISS",
      "name": "Kate Rubins"
    },
    {
      "craft": "ISS",
      "name": "Soichi Noguchi"
    }
  ]
}
```

```
}  
}  
}
```

You may also like to try out the bash template. You can search for templates in the store with `faas-cli template store list` and then fetch it to create a new function.

NAME	SOURCE	DESCRIPTION
bash-streaming	openfaas-incubator	Bash Streaming template

```
faas-cli template store pull bash-streaming  
faas-cli new --lang bash-streaming bash-script
```

Chapter 5

Monitoring invocations

You can monitor invocations using the built-in Prometheus metrics, or by deploying Grafana as an additional component in the faasd stack.

What's the difference between the two? Prometheus is a time-series database which has a basic UI built into it for visualising queries and Grafana is a UI tool for creating dashboards from many different data-sources like Prometheus and InfluxDB.

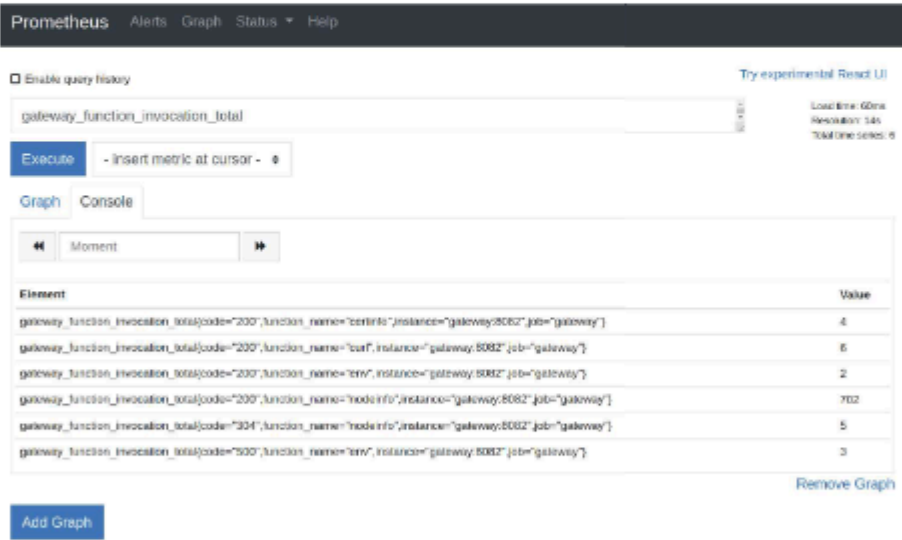


Figure 5.1: Prometheus UI

The Prometheus UI showing the query executor

Create an SSH tunnel to faasd and forward the Prometheus UI back to your computer:

```
export IP="faasd-ip"
ssh -L 9090:127.0.0.1:9090 ubuntu@IP
```

Note that the user may vary depending on where you are hosting your faasd VM.

Now navigate to <http://localhost:9090>

You can view the various metrics using [PromQL](#).

View the rate of new invocations started

```
rate(gateway_invocation_started [1m])
```

View the services which are scaled to zero or one replica:

```
gateway_service_count
```

The OpenFaaS REST API can be monitored through:

```
http_requests_total
```

```
http_request_duration_seconds
```

You'll be able to query Prometheus in the browser and see how long executions are taking, and whether any are failing with non 200 HTTP codes.

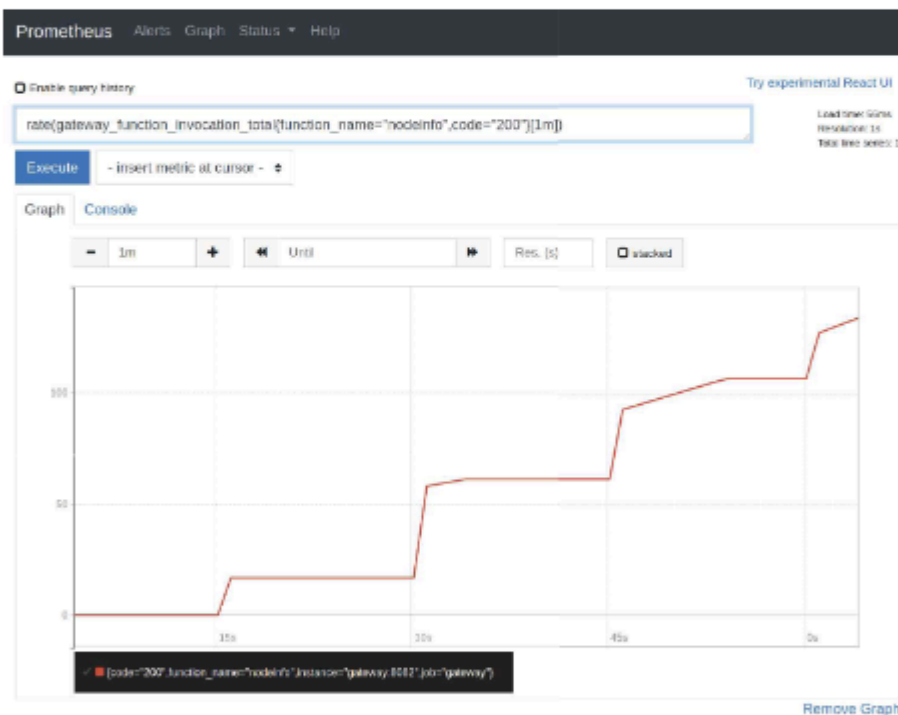


Figure 5.2: Prometheus graph

The Prometheus UI showing a graph of the execution rate going up under a load test

See also: [built-in metrics](#)

In the "Additional containers and services" section you'll learn how to deploy a Grafana dashboard to monitor the data from Prometheus.